

Peak process GPU memory

Measured context/batch groups at cache ratios 0.5%, 5%, and 10%

